

# Hands on with Generative Adversarial Networks (GANs): Making some pretty pictures

Nabeel Seedat

shutterstock®



## SUMMARY

1. Intuition behind GANs
2. How GANs work
3. Code Hands-on
  - GAN for making new pictures of clothes
  - Tips & Tricks to make GANs work
4. Where to from here?

<http://tinyurl.com/pygotham-GANS>

# FACE GENERATION




2014



2015



2016



2017



2018

[Goodfellow](#)

# IMAGE-TO-IMAGE GENERATION

Labels to Street Scene




input



output

Aerial to Map




input



output

Labels to Facade



input



output

BW to Color



input



output

Day to Night



input



output

Edges to Photo



input



output

[CGAN \(Isola et al\)](#)



Monet  $\leftrightarrow$  Photos



Monet  $\rightarrow$  photo




photo  $\rightarrow$  Monet

[CycleGAN \(Xu et al\)](#)

Zebras  $\leftrightarrow$  Horses



zebra  $\rightarrow$  horse



horse  $\rightarrow$  zebra

# TEXT-TO-IMAGE GENERATION

this bird is red with white and has a very short beak



[AttnGAN \(Xu et al\)](#)

# SOME DEEPPAKES TOO



[Derpfakes](#)

# GENERATIVE ADVERSARIAL NETWORKS (GANS) MAKE THIS POSSIBLE!

Goal: Generate new data!

- Generative  
*Learn a generative model*
- Adversarial  
*Model trained in an adversarial setting*
- Networks  
*Use Deep Neural Networks*



# Generative?

- Most deep learning is discriminative

Given image ( $x$ ) - predict label ( $y$ )

$$p(y|x)$$

- Generative - don't care about labels

*Learn how the data ( $x$ ) was generated =  $p(x)$*

*If we know  $p(x)$ , we can generate new images!*

- *Deeper look in the notebook:*  
BUT note unsupervised is not always generative (e.g. K-Means clustering)


# Adversarial Intuition

Generator


Discriminator



FAKE



REAL




Real  
Data


Image Sources: Shutterstock

# Intuition behind GANs

Generator




Discriminator



FAKE



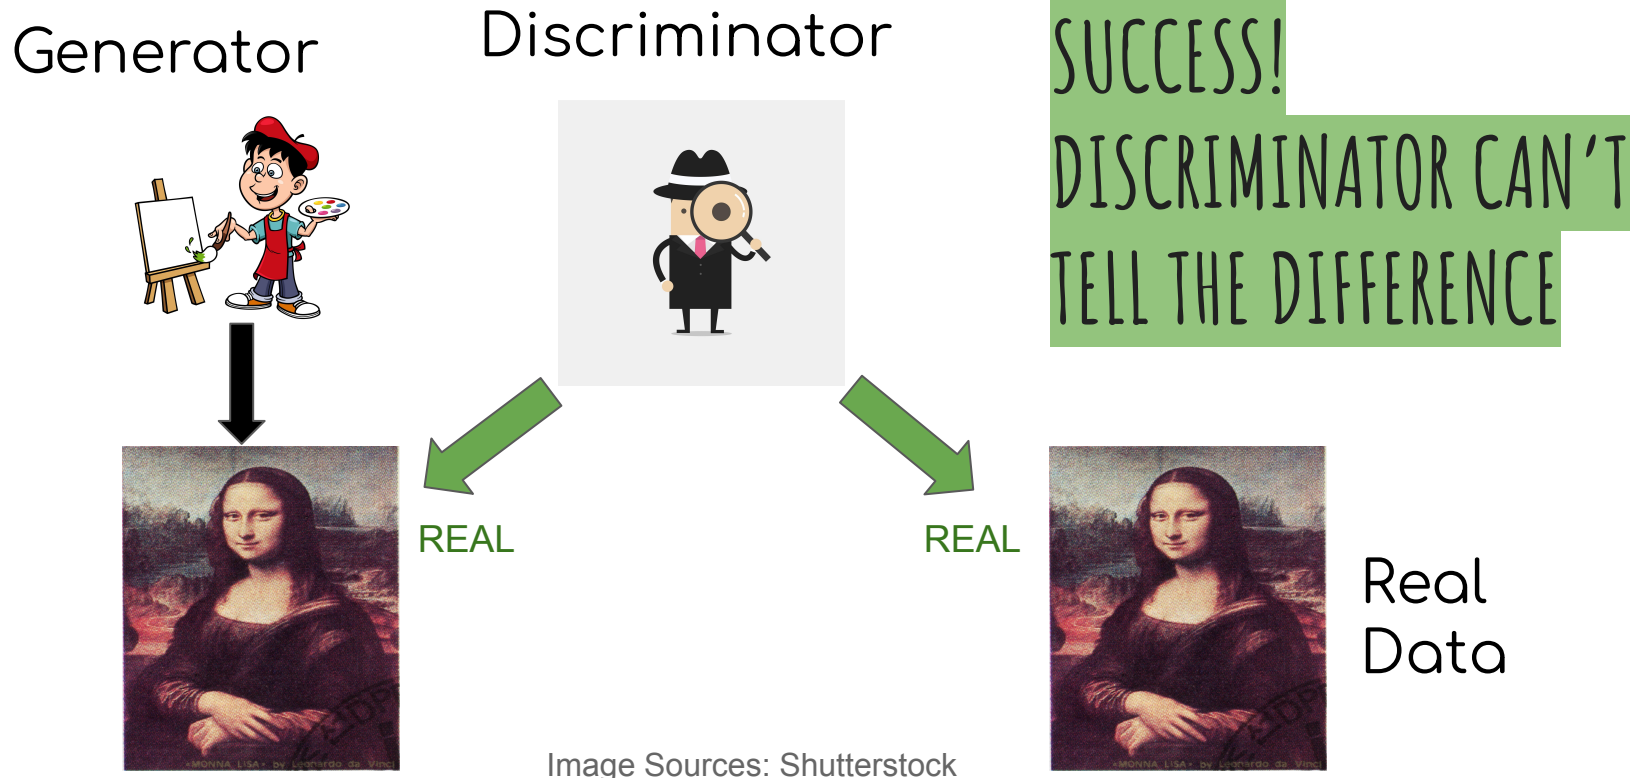
REAL



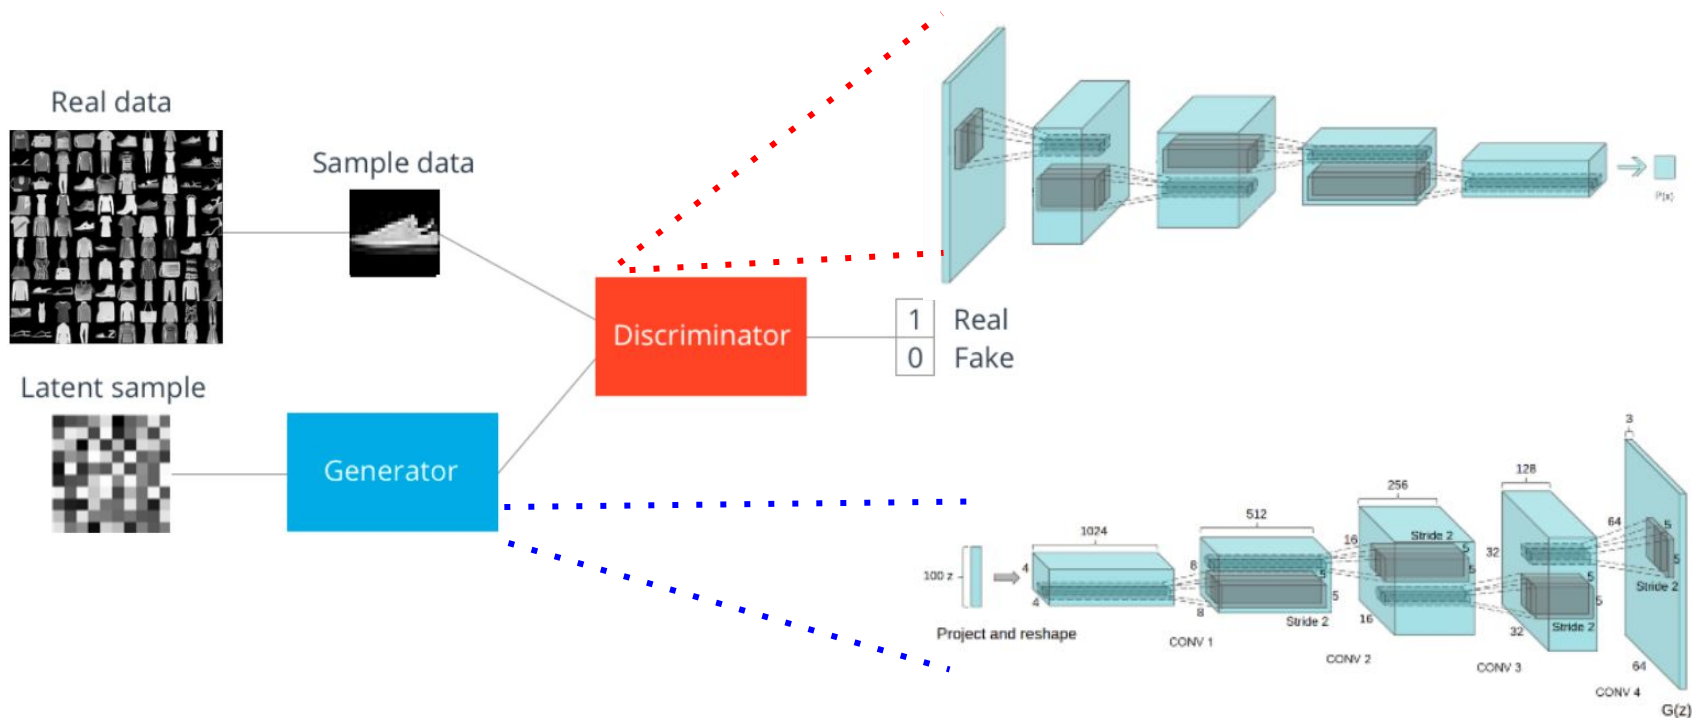
Real  
Data

Image Sources: Shutterstock

# Intuition behind GANs



# How GANs work



*This uses transpose convolution*

DCGAN: Radford et al



# LOSS FUNCTION - MINIMAX

Generator tries to minimize the reward of the discriminator (by fooling it)

Discriminator tries to maximize reward (predicting correctly)

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Log probability of D predicting the real data is actually real


Log probability of D predicting the data from G is NOT real

*Generator plays no role so during generator training we don't need this*

*This has problems of vanishing gradients - BUT let's ignore it & chat at the end*

# GAN Hands On!

- Code to follow along can be found on Github
- It's also a self-contained tutorial in a notebook linked to Google Colab (with a free GPU!)




Github link:

<http://tinyurl.com/pygotham-GANS>

# DATA PRE-PROCESSING

```
images = np.expand_dims(train, axis=3) # add a channel dim
images = images.astype('float32') # convert to float32
images = (images - 127.5) / 127.5 # normalize the images
batch_size = 256
dataset =
tf.data.Dataset.from_tensor_slices(images).shuffle(batch_size*10).batch(batch_
size)
```

# DISCRIMINATOR



```
def discriminator():
```

```
    discriminator = tf.keras.Sequential(name="discriminator")
```

```
    discriminator.add(tf.keras.layers.Conv2D(32, kernel_size=3, strides=2, padding="same", input_shape=(28, 28, 1)))
```

```
    discriminator.add(tf.keras.layers.LeakyReLU(alpha=0.2))
```

```
    discriminator.add(tf.keras.layers.Dropout(0.25))
```

```
    discriminator.add(tf.keras.layers.Conv2D(64, kernel_size=3, strides=2, padding="same"))
```

```
    discriminator.add(tf.keras.layers.BatchNormalization(momentum=0.8))
```

```
    discriminator.add(tf.keras.layers.LeakyReLU(alpha=0.2))
```

```
    discriminator.add(tf.keras.layers.Dropout(0.25))
```

```
    discriminator.add(tf.keras.layers.Conv2D(128, kernel_size=3, strides=2, padding="same"))
```

```
    discriminator.add(tf.keras.layers.BatchNormalization(momentum=0.8))
```

```
    discriminator.add(tf.keras.layers.LeakyReLU(alpha=0.2))
```

```
    discriminator.add(tf.keras.layers.Dropout(0.25))
```

```
    discriminator.add(tf.keras.layers.Conv2D(256, kernel_size=3, strides=1, padding="same"))
```

```
    discriminator.add(tf.keras.layers.BatchNormalization(momentum=0.8))
```

```
    discriminator.add(tf.keras.layers.LeakyReLU(alpha=0.2))
```


```
    discriminator.add(tf.keras.layers.Dropout(0.25))
```

```
    discriminator.add(tf.keras.layers.Flatten())
```

```
    discriminator.add(tf.keras.layers.Dense(1, activation='sigmoid'))
```

```
    return discriminator
```

# GENERATOR



```
def generator():
    generator = tf.keras.Sequential(name="generator")
```

```
    generator.add(tf.keras.layers.Dense(7 * 7 * 128, activation="relu", input_dim=100))
    generator.add(tf.keras.layers.Reshape([7, 7, 128]))
```

```
    generator.add(tf.keras.layers.UpSampling2D())
    generator.add(tf.keras.layers.Conv2D(128, kernel_size=3, padding="same"))
    generator.add(tf.keras.layers.BatchNormalization(momentum=0.8))
    generator.add(tf.keras.layers.Activation("relu")) ★
```

} *Transpose convolution*

```
    generator.add(tf.keras.layers.UpSampling2D())
    generator.add(tf.keras.layers.Conv2D(64, kernel_size=3, padding="same"))
    generator.add(tf.keras.layers.BatchNormalization(momentum=0.8))
    generator.add(tf.keras.layers.Activation("relu")) ★
```

} *Transpose convolution*

```
    generator.add(tf.keras.layers.Conv2D(1, kernel_size=3, padding="same"))
    generator.add(tf.keras.layers.Activation("tanh")) ★
```

*This outputs data= [-1,1]*

```
    return generator
```



# LOSS FUNCTIONS

```
# define the loss function for the discriminator  
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=False)
```

```
def discriminator_loss(real_images, fake_images):
```

```
    real_labels=tf.ones_like(real_images)  
    real_loss = cross_entropy(real_labels, real_images)
```

Loss: real  
images

```
    fake_labels=tf.zeros_like(fake_images)  
    fake_loss = cross_entropy(fake_labels, fake_images)
```

Loss: fake  
images

```
    total_loss = real_loss + fake_loss  
    return total_loss
```

Discriminator Loss

```
def generator_loss(fake_images):  
    fake_labels=tf.zeros_like(fake_images)  
    return cross_entropy(fake_labels, fake_images)
```

Generator Loss

# GAN TRAINING

```
@tf.function  
def train_step(images):  
    noise = tf.random.normal([batch_size, 100])
```

1. Generate random noise

Latent sample



# GAN TRAINING

```
@tf.function
def train_step(images):
    noise = tf.random.normal([batch_size, 100])
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)
```

2. Generate fake image using noise

Latent sample



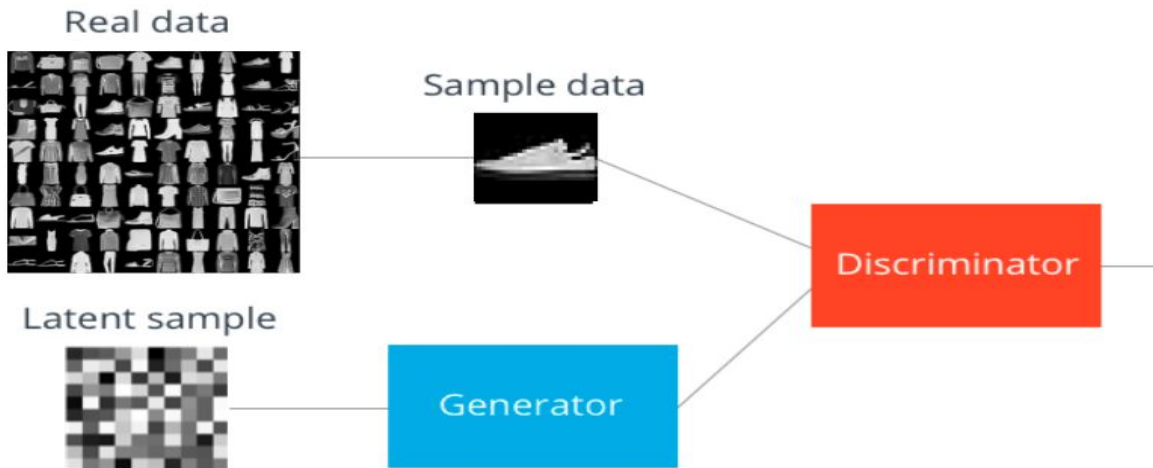
Generator



# GAN TRAINING

```
@tf.function
def train_step(images):
    noise = tf.random.normal([batch_size, 100])
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)
        real_images = discriminator(images, training=True)
        fake_images = discriminator(generated_images, training=True)
```

3. Pass real & fake images through the discriminator



# GAN TRAINING

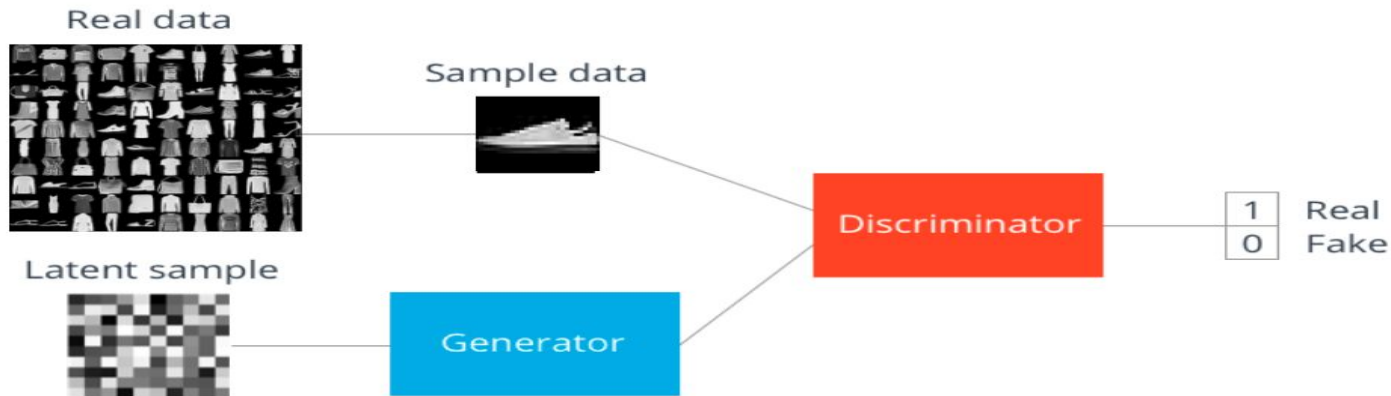
```
@tf.function
def train_step(images):
    noise = tf.random.normal([batch_size, 100])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_images = discriminator(images, training=True)
        fake_images = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_images)
        disc_loss = discriminator_loss(real_images, fake_images)
```

4. Compute Losses for both





# GAN TRAINING

```
@tf.function
def train_step(images):
    noise = tf.random.normal([batch_size, 100])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

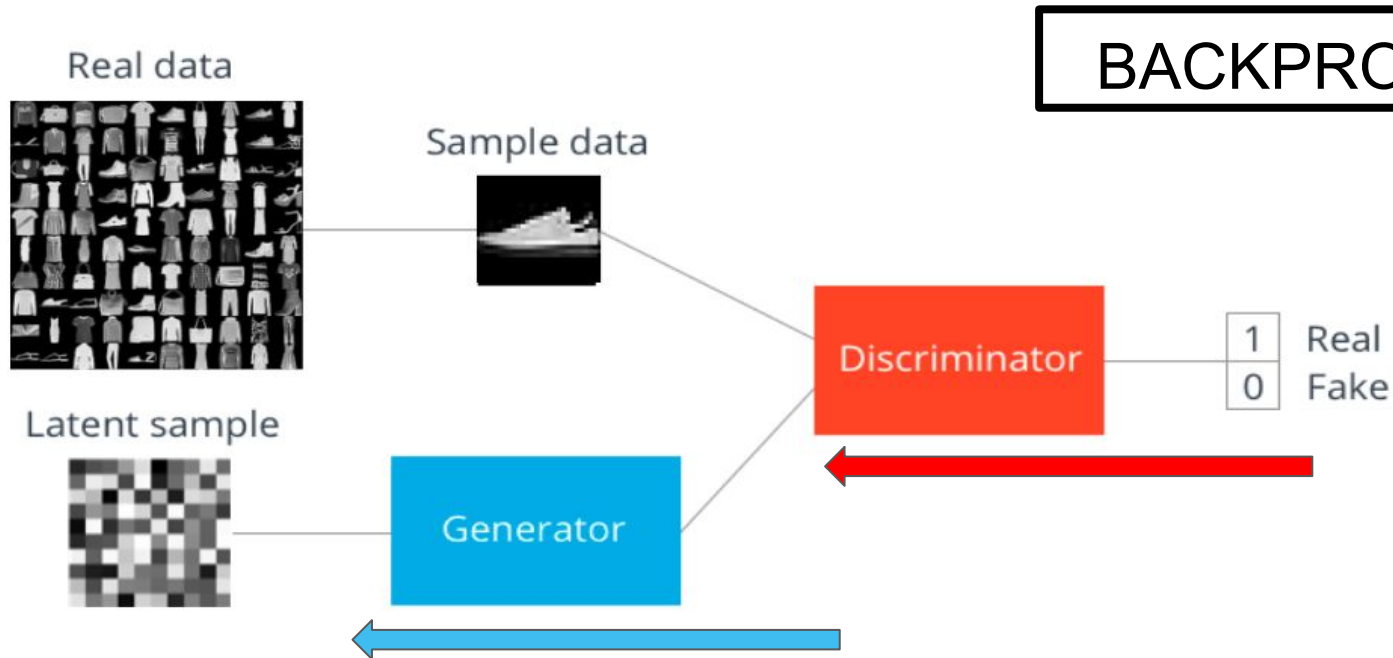
        real_images = discriminator(images, training=True)
        fake_images = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_images)
        disc_loss = discriminator_loss(real_images, fake_images)

    gen_loss_mean(gen_loss)
    disc_loss_mean(disc_loss)
```

5. Get avg loss for the epoch

# GAN TRAINING



# GAN TRAINING

```
@tf.function
def train_step(images):
    noise = tf.random.normal([batch_size, 100])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_images = discriminator(images, training=True)
        fake_images = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_images)
        disc_loss = discriminator_loss(real_images, fake_images)

    gen_loss_mean(gen_loss)
    disc_loss_mean(disc_loss)

    gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)
    gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)
```

6. Compute gradient.  
(Autodiff for backprop)

# GAN TRAINING

```
@tf.function
def train_step(images):
    noise = tf.random.normal([batch_size, 100])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_images = discriminator(images, training=True)
        fake_images = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_images)
        disc_loss = discriminator_loss(real_images, fake_images)


    gen_loss_mean(gen_loss)
    disc_loss_mean(disc_loss)

    gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)
    gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)

    generator_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
    discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator,
discriminator.trainable_variables))
```

7. Apply optimizer. Gradient step

# RESULTS





# Where to from here?

- Deeper problems + solutions
- Beyond the vanilla type GAN for more fun examples

# DEEPER PROBLEMS AND SOLUTIONS

- *Not converging!*: Discriminator is too good so generator fails

*Solution: Add noise to discriminator input ([Arjovsky & Bottou](#))*

- *Vanishing gradients*: Discriminator is too good so generator fails

*Solution: Use Wasserstein Loss ([link](#))*

- *Mode collapse*: Generator produces the same examples

*Solution: Wasserstein Loss or Unrolled GAN ([Metz et al](#))*

# FUN EXAMPLES - GAN ZOO

- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks

- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

<https://github.com/hindupuravinash/the-gan-zoo>

THANK YOU!

Questions?